

50

PYTHON EXERCISES



NUMPY
PANDAS

Daniel Christofis

50 PYTHON EXERCISES

NUMPY & PANDAS

A Practical Guide - Mastering DataFrames, Data
Manipulation and Creating Graphs

Daniel Christofis

CONTENTS

[Title Page](#)

[A quick Note](#)

[Introduction](#)

[Installing Required Libraries](#)

[Project 1: Introduction to Pandas and NumPy Basics](#)

[Project 2: Data Filtering & Sorting](#)

[Project 3: Data Grouping & Aggregation](#)

[Project 4: Data Cleaning & Exploration](#)

[Project 5: Merging, Joining & Concatenation](#)

[Project 6: Visualization & Descriptive Statistics](#)

[Project 7: Advanced Pandas Functions & Techniques](#)

[Project 8: Feature Engineering & Selection Project](#)

[Project 1 Solutions](#)

[Project 2 Solutions](#)

[Project 3 Solutions](#)

[Project 4 Solutions](#)

[Project 5 Solutions](#)

[Project 6 Solutions](#)

[Project 7 Solutions](#)

[Project 8 Solutions](#)

A QUICK NOTE

We hope you are enjoying your recent purchase from Amazon! We wanted to take a moment to kindly remind you about the importance of leaving a review for the product you bought. Your feedback is incredibly valuable to us and other potential buyers.

By sharing your thoughts and experiences with the product, you help others make informed decisions. Your review can be a guiding light for fellow customers, ensuring they choose the best products for their needs.


We would greatly appreciate it if you could spare a few moments to write an honest review. Your feedback helps us improve our products and services, making it a win-win situation for everyone.


Thank you for being a valued customer and for considering writing a review. We truly appreciate your support.

Introducing my personal blog. Stay connected!

<https://premiumbookself.wordpress.com/>







INTRODUCTION

NumPy and Pandas emerge as the champions of the Python data environment in the enormous area of data research. These libraries allow not only professionals to easily manage and alter data, but also newcomers to dig deep into the field of data analysis with ease and intuitively. "50 Python Exercises: NumPy and Pandas" has been carefully designed to give you a taste of real-world issues, allowing you to flex your problem-solving muscles and hone your coding skills.

This book provides a wide range of activities organised into eight projects. Each project focuses on a distinct topic, bringing you from the fundamentals of data manipulation to the depths of feature engineering.

Project 1: Introduction to Pandas and NumPy Basics

Grasp the foundations of data handling. Dive into tasks ranging from loading datasets to basic computations, setting the stage for more advanced projects.

Project 2: Data Filtering & Sorting

Learn how to sift through data and present it in a way that makes sense. This section emphasizes the importance of structuring and organizing data for better analysis.

Project 3: Data Grouping & Aggregation

Delve deeper into the data, dissecting it into meaningful groups and drawing aggregated insights that can drive business decisions or power data models.

Project 4: Data Cleaning & Exploration

Every data scientist knows that raw data is often messy. Here, we take on the vital tasks of cleaning, transforming, and exploring datasets to prepare them for further analysis or modeling.

Project 5: Merging, Joining & Concatenation

Data doesn't always come in a single, neat package. Discover how to combine datasets, be it stacking them up or joining them based on common attributes.

Project 6: Visualization & Descriptive Statistics

Pictures often speak louder than numbers. Dive into the art and science of visualizing data, alongside wielding descriptive statistics to capture the essence of your dataset.

Project 7: Advanced Pandas Functions & Techniques

Step into the realm of advanced data manipulation. This section delves into some of the more intricate functionalities that Pandas offers, which can immensely streamline your data processing workflow.

Project 8: Feature Engineering & Selection

In the world of machine learning and predictive modeling, the features you provide to your model are paramount. Master the art of crafting, selecting, and transforming features to boost the performance of subsequent models.

INSTALLING REQUIRED LIBRARIES

Before diving into the exercises, it's essential to ensure that you have all the required libraries installed. This book primarily leverages the power of NumPy and Pandas, along with a few additional libraries for specific tasks.

To install the necessary libraries, you can use Python's package manager, pip. Simply open your terminal or command prompt and enter the following command:



```
pip install numpy pandas seaborn matplotlib scikit-learn statsmodels
```

Once the installation is complete, you're all set to embark on the exciting journey of "50 Python Exercises: NumPy and Pandas". **Happy coding!**

PROJECT 1: INTRODUCTION TO PANDAS AND NUMPY BASICS

Import the 'NumPy' and 'Pandas' libraries, and the 'Titanic' dataset from seaborn.

```
import seaborn as sns  
import numpy as np  
import pandas as pd  
  
titanic = sns.load_dataset('titanic')
```

- 1-1. Load the Titanic dataset using Seaborn and display the first 5 rows.
- 1-2. Display the last 5 rows of the dataset.
- 1-3. Get the data types of each column.
- 1-4. Convert the 'sex' column to a category type.
- 1-5. Compute the sum of the 'fare' column.
- 1-6. Get the mean and standard deviation of the 'age' column.
- 1-7. Create a NumPy array from the 'age' column of the Titanic DataFrame.

PROJECT 2: DATA FILTERING & SORTING

Import the 'NumPy' and 'Pandas' libraries, and the 'Titanic' dataset from seaborn.

```
import seaborn as sns  
import numpy as np  
import pandas as pd
```

```
titanic = sns.load_dataset('titanic')
```

- 2-1. Filter out rows where the 'age' column's value is missing.
- 2-2. Retrieve rows where the age is greater than 50.
- 2-3. Filter data based on multiple conditions (e.g., $\text{age} > 30$ and 'sex' is 'female').
- 2-4. Sort the dataset based on the 'fare' column in ascending order.
- 2-5. Sort the dataset by multiple columns (e.g., first by 'class' and then by 'age').
- 2-6. Reset the index of the DataFrame after sorting.
- 2-7. Find and display rows where 'embarked' column value is 'C' (Cherbourg) and 'fare' is within the 90th percentile.

PROJECT 3: DATA GROUPING & AGGREGATION

Import the 'Pandas' library, and the 'iris' dataset from seaborn.

```
import pandas as pd  
import seaborn as sns
```

```
iris = sns.load_dataset('iris')
```

3-1: Display the first 5 rows of the dataset.

3-2: Get a description (count, mean, std, min, max) of each numeric column in the dataset

3-3: Group the dataset by the 'species' column and get the mean value of each column for each species.

3-4: Which species has the highest average sepal length?

3-5: Filter the dataset for the 'setosa' species. How many unique values are there in the 'petal_width' column for this species, and what is the count for each unique value?

3-6: Identify the species with the maximum petal width. Additionally, find its corresponding average sepal width.

3-7: Create a pivot table with species as rows. The table should display the average petal length and sepal length for each species.

PROJECT 4: DATA CLEANING & EXPLORATION

Dataset:

```
import seaborn as sns
import pandas as pd

# Loading the mpg dataset from seaborn
mpg = sns.load_dataset('mpg')
```

4-1: Display the last 7 rows of the dataset.

4-2: How many missing values are there in the 'horsepower' column? Replace them with the median value of that column.

4-3: Filter the dataset for cars with an 'origin' from 'usa' and 'cylinders' equal to 4. How many such cars are there in the dataset?

4-4: Identify the make and model of the car with the highest 'mpg'. What is its mpg?

4-5: Create a new column 'displacement_to_power' as the ratio of 'displacement' to 'horsepower'. Which car make and model has the highest 'displacement_to_power' ratio?

4-6: Group the dataset by 'origin' and calculate the average 'mpg' for each group. Which 'origin' has the highest average mpg?

4-7: Filter cars with an 'mpg' lower than the 10th percentile. How many cars remain in the dataset post this operation?

PROJECT 5: MERGING, JOINING & CONCATENATION

Import both datasets:

```
import seaborn as sns  
import pandas as pd
```

```
# Loading the datasets  
tips = sns.load_dataset('tips')  
flights = sns.load_dataset('flights')
```

5-1: Concatenate 'tips' and 'flights' DataFrames vertically (keep in mind the datasets are different; this is more for practice than utility).

5-2: Merge 'tips' and a DataFrame that has an additional column, say "discount", based on a common column "total_bill".

5-3: Perform a left join between the 'tips' DataFrame and another DataFrame containing "total_bill" and a new column "feedback".

5-4: Perform an outer join between 'tips' and a DataFrame with columns "total_bill" and "restaurant_review".

5-5: Concatenate 'tips' and 'flights' DataFrames horizontally.

5-6: Handle overlapping columns when merging two DataFrames: 'tips' and a DataFrame containing "total_bill", "tip", and "dining_experience".

5-7: Merge 'tips' and a DataFrame using a multi-index based on columns "total_bill" and "tip".

PROJECT 6: VISUALIZATION & DESCRIPTIVE STATISTICS

Dataset and extra libraries used:

```
import seaborn as sns  
import pandas as pd  
import matplotlib.pyplot as plt  
import scipy.stats as stats  
  
tips = sns.load_dataset('tips')
```

- 6-1. Generate a histogram for the 'total_bill' column using Pandas.
- 6-2. Compute and display the correlation matrix for the entire tips dataset.
- 6-3. Generate a scatter plot between the 'total_bill' and 'tip' columns using Pandas.
- 6-4. Calculate and display the skewness and kurtosis for the 'total_bill' column.
- 6-5. Use the value_counts method to display the frequency of unique values in the 'day' column. Subsequently, visualize it as a bar chart.
- 6-6. Generate a box plot for the 'total_bill' column to visualize outliers.
- 6-7. Compute the IQR (Interquartile Range) for the 'total_bill' column.

PROJECT 7: ADVANCED PANDAS FUNCTIONS & TECHNIQUES

Dataset:

```
import seaborn as sns
import numpy as np
import pandas as pd

diamonds = sns.load_dataset('diamonds')
```

7-1. Use the query method to filter rows from the diamonds dataset where the cut is "Ideal" and the price is greater than \$500.

7-2. Extract rows from the diamonds dataset where the color column matches the regex pattern "^EF\$".

7-3. Set the 'color' column of the diamonds dataset as the index, and then reset it back to a column.

7-4. Use the eval function to compute the price divided by carat and store the result in a new column named "price_per_carat".

7-5. Create a new column named 'high_price' in the diamonds dataset, which should be True if the price is greater than \$1000 and False otherwise. Use the np.where function for this.

7-6. Find and drop columns from the diamonds dataset with a variance below a threshold of 1.

7-7. Use the melt function to reshape the diamonds dataset, with 'color' and 'cut' as ID variables.

PROJECT 8: FEATURE ENGINEERING & SELECTION PROJECT (Optional)

Dataset and extra libraries used:

```
import seaborn as sns
import pandas as pd
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.decomposition import PCA
from statsmodels.stats.outliers_influence import
    variance_inflation_factor
titanic = sns.load_dataset('titanic')
```

8-1. Based on existing columns in the Titanic dataset, namely 'sibsp' which denotes the number of siblings/spouses aboard, and 'parch' which signifies the number of parents/children aboard, can you create a new column called 'family_size' to represent the total family size of each passenger on board?

8-2. The 'class' column in the Titanic dataset refers to the passenger class and is a categorical column. Can you transform this column using one-hot encoding, ensuring that you drop the first category to avoid multicollinearity?

8-3. Age is an essential continuous variable in the Titanic dataset. Can you normalize the 'age' column using the z-score normalization technique? This helps to express ages in terms of their deviation from the mean age.

8-4. Dimensionality reduction can help in reducing noise and improving computational efficiency. Using PCA (Principal Component Analysis), can you reduce the dimensions of the dataset? For this task, consider only the 'age' and 'fare' columns.

8-5. Multicollinearity can be an issue when building predictive models. For the Titanic dataset, can you compute the Variance Inflation Factor (VIF) for the 'age', 'fare', and the one-hot encoded 'class' columns? Features with a VIF greater than 10 are typically considered to have high multicollinearity.

8-6. For better interpretability and categorization, can you bin the 'age' column of the Titanic dataset into custom groups: 'child' (0-18 years), 'young_adult' (18-35 years), 'adult' (35-60 years), and 'senior' (60+ years)?

8-7. Target mean encoding, also known as likelihood encoding, is a way to encode categorical columns based on the mean of the target variable. Using this technique, can you encode the 'embark_town' column in the Titanic dataset with respect to the 'survived' column?

PROJECT 1 SOLUTIONS

```
import seaborn as sns
import numpy as np
import pandas as pd

# Load the dataset
titanic = sns.load_dataset('titanic')

# 1-1. Load the Titanic dataset using Seaborn and display the first 5 rows.
print("#1-1:\n", titanic.head())

# 1-2. Display the last 5 rows of the dataset.
print("#1-2:\n", titanic.tail())

# 1-3. Get the data types of each column.
print("#1-3:\n", titanic.dtypes)

# 1-4. Convert the 'sex' column to a category type.
titanic['sex'] = titanic['sex'].astype('category')
print("#1-4:\n", titanic['sex'].dtypes)

# 1-5. Compute the sum of the 'fare' column.
fare_sum = titanic['fare'].sum()
print("#1-5: Sum of Fare Column:", fare_sum)

# 1-6. Get the mean and standard deviation of the 'age' column.
age_mean = titanic['age'].mean()
age_std = titanic['age'].std()
print("#1-6:\nMean Age:", age_mean, "\nStandard Deviation of Age:",
age_std)

# 1-7. Create a NumPy array from the 'age' column of the Titanic DataFrame.
age_array = titanic['age'].to_numpy()
print("#1-7:\n", age_array)
```

Project 1 Output

#1-1:

```
survived pclass sex age sibsp parch ...
0      0     3 male 22.0   1   0
1      1     1 female 38.0   1   0
2      1     3 female 26.0   0   0
3      1     1 female 35.0   1   0
4      0     3 male 35.0   0   0
```

#1-2:

```
survived pclass sex age sibsp parch ...
886      0     2 male 27.0   0   0
887      1     1 female 19.0   0   0
888      0     3 female NaN   1   2
889      1     1 male 26.0   0   0
890      0     3 male 32.0   0   0
```

#1-3:

```
survived      int64
pclass       int64
sex         object
age        float64
sibsp       int64
parch       int64
fare        float64
embarked    object
class       category
who         object
adult_male   bool
deck       category
embark_town object
alive       object
alone       bool
dtype: object
```

#1-4:

category

#1-5: Sum of Fare Column: 28693.9493

#1-6:

Mean Age: 29.69911764705882

Standard Deviation of Age: 14.526497332334044

#1-7:

[22. 38. 26. ... nan 26. 32.]

PROJECT 2 SOLUTIONS

```
import seaborn as sns
import numpy as np
import pandas as pd

titanic = sns.load_dataset('titanic')

# 2-1. Filter out rows where the 'age' column's value is missing.
age_not_missing = titanic.dropna(subset=['age'])
print("#2-1", age_not_missing.head())

# 2-2. Retrieve rows where the age is greater than 50.
age_above_50 = titanic[titanic['age'] > 50]
print("#2-2", age_above_50.head())

# 2-3. Filter data based on multiple conditions (e.g., age > 30 and 'sex' is 'female').
females_above_30 = titanic[(titanic['age'] > 30) & (titanic['sex'] == 'female')]
print("#2-3", females_above_30.head())

# 2-4. Sort the dataset based on the 'fare' column in ascending order.
sorted_by_fare = titanic.sort_values(by='fare')
print("#2-4", sorted_by_fare.head())

# 2-5. Sort the dataset by multiple columns (e.g., first by 'class' and then by 'age').
sorted_by_class_age = titanic.sort_values(by=['class', 'age'])
print("#2-5", sorted_by_class_age.head())

# 2-6. Reset the index of the DataFrame after sorting.
reset_index_df = sorted_by_class_age.reset_index(drop=True)
print("#2-6", reset_index_df.head())
```

```
# 2-7. Find and display rows where 'embarked' column value is
'C' (Cherbourg) and 'fare' is within the 90th percentile.
cherbourg_high_fare = titanic[(titanic['embarked'] == 'C') &
(titanic['fare'] > titanic['fare'].quantile(0.90))]
print("#2-7", cherbourg_high_fare)
```

Project 2 Output

#2-1:

	survived	pclass	sex	age	sibsp	parch	...
0	0	3	male	22.0	1	0	
1	1	1	female	38.0	1	0	
2	1	3	female	26.0	0	0	
3	1	1	female	35.0	1	0	
4	0	3	male	35.0	0	0	

#2-2:

	survived	pclass	sex	age	sibsp	parch	...
6	0	1	male	54.0	0	0	
11	1	1	female	58.0	0	0	
15	1	2	female	55.0	0	0	

#2-3:

	survived	pclass	sex	age	sibsp	parch	fare	...
1	1	1	female	38.0	1	0	71.2833	
3	1	1	female	35.0	1	0	53.1000	
11	1	1	female	58.0	0	0	26.5500	

#2-4:

	survived	pclass	sex	age	sibsp	parch	fare	...
179	0	3	male	NaN	0	0	0.0	
263	0	1	male	NaN	0	0	0.0	
302	0	3	male	NaN	0	0	0.0	

#2-5:

	survived	pclass	sex	age	...	class	who	...
--	----------	--------	-----	-----	-----	-------	-----	-----

```
20     0     2   male  35.0      Second   man
21     1     2   male  34.0      Second   man
33     0     2   male  66.0      Second   man
```

#2-6:

```
  survived pclass   sex   age sibsp parch ...
0         0     2   male  35.0     0     0
1         1     2   male  34.0     0     0
2         0     2   male  66.0     0     0
```

#2-7:

```
survived... fare embarked ... deck ... embark_town alive alone
31     1    146.5208      C  Cherbourg yes False
195    1    146.5208      C  Cherbourg yes True
305    1    151.5500      C  Cherbourg yes False
```

PROJECT 3 SOLUTIONS

```
import pandas as pd  
import seaborn as sns
```

```
iris = sns.load_dataset('iris')
```

3-1: Display the first 5 rows of the dataset.

```
print(iris.head())
```

3-2: Get a description (count, mean, std, min, max) of each numeric column.

```
print(iris.describe())
```

3-3: Group the dataset by 'species' and get the mean value of each column per species.

```
grouped_by_species = iris.groupby('species').mean()  
print(grouped_by_species)
```

3-4: Find the species that has the highest average sepal length.

```
max_sepal_length_species =  
grouped_by_species['sepal_length'].idxmax()  
print("Species with the highest average sepal length:",  
max_sepal_length_species)
```

3-5: Filter the data for setosa species and get the count of each unique value in the 'petal_width' column.

```
setosa_petal_width_count =  
iris[iris['species'] == 'setosa']['petal_width'].value_counts()  
print(setosa_petal_width_count)
```

3-6: Using aggregation, find the species with the maximum petal width and its corresponding average sepal width.

```
agg_data = iris.groupby('species').agg({'petal_width': 'max',  
'sepal_width': 'mean'})
```

```
species_max_petal_width = agg_data['petal_width'].idxmax()
avg_sepal_width = agg_data.loc[species_max_petal_width,
                                'sepal_width']
print(f"Species with maximum petal width:
{species_max_petal_width}, with avg sepal width:
{avg_sepal_width}")
```

3-7: Create a pivot table with species as rows, and average petal length and sepal length as values.

```
pivot_table = iris.pivot_table(index='species', values=['petal_length',
                                                       'sepal_length'], aggfunc='mean')
print(pivot_table)
```

Project 3 Output

#3-1:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

#3-2:

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000

...

(Note: I've abbreviated the describe output for brevity.)

#3-3:

species	sepal_length	sepal_width	petal_length	petal_width
---------	--------------	-------------	--------------	-------------

setosa	5.006	3.428	1.462	0.246
versicolor	5.936	2.770	4.260	1.326
virginica	6.588	2.974	5.552	2.026

#3-4:

Species with the highest average sepal length: virginica

#3-5:

0.2	28
0.4	7
0.3	7
0.1	6
0.5	1
0.6	1

Name: petal_width, dtype: int64

#3-6:

Species with maximum petal width: virginica, with avg sepal width: 2.974

#3-7:

	petal_length	sepal_length
species		
setosa	1.462	5.006
versicolor	4.260	5.936
virginica	5.552	6.588

PROJECT 4 SOLUTIONS

```
import seaborn as sns
import pandas as pd

# Loading the mpg dataset from seaborn
mpg = sns.load_dataset('mpg')

# 4-1: Display the last 7 rows of the dataset.
print("#4-1:")
print(mpg.tail(7))

# 4-2: Find missing values in 'horsepower' and replace with median.
missing_values_count = mpg['horsepower'].isna().sum()
mpg['horsepower'].fillna(mpg['horsepower'].median(), inplace=True)
print("\n#4-2:")
print(f"Number of missing values replaced in 'horsepower':\n{missing_values_count}")

# 4-3: Filter the data for 'usa' origin and 4 cylinders.
filtered_cars = mpg[(mpg['origin'] == 'usa') & (mpg['cylinders'] == 4)]
print("\n#4-3:")
print(f"Number of cars from 'usa' with 4 cylinders:\n{len(filtered_cars)}")

# 4-4: Identify the car with the highest mpg.
highest_mpg_row = mpg[mpg['mpg'] == mpg['mpg'].max()]
print("\n#4-4:")
print(f"""\nCar with the highest mpg:
{highest_mpg_row['name'].values[0]} - MPG:
{highest_mpg_row['mpg'].values[0]}""")

# 4-5: Create 'displacement_to_power' and find the car with the highest ratio.
mpg['displacement_to_power'] = mpg['displacement'] /
```

```

mpg['horsepower']
highest_ratio_row = mpg[mpg['displacement_to_power'] ==
mpg['displacement_to_power'].max()]
print("#4-5:")
print(f"Car with the highest 'displacement_to_power' ratio:
{highest_ratio_row['name'].values[0]} - Ratio:
{highest_ratio_row['displacement_to_power'].values[0]}")

# 4-6: Group by 'origin' and calculate average mpg.
avg_mpg_by_origin = mpg.groupby('origin')['mpg'].mean()
highest_avg_mpg_origin = avg_mpg_by_origin.idxmax()
print("\n#4-6:")
print(f"Origin with the highest average mpg:
{highest_avg_mpg_origin} - Average MPG:
{avg_mpg_by_origin.max():.2f}")

# 4-7: Filter cars with mpg lower than 10th percentile.
mpg_10th_percentile = mpg['mpg'].quantile(0.10)
filtered_mpg = mpg[mpg['mpg'] > mpg_10th_percentile]
print("\n#4-7:")
print(f"Number of cars remaining after removing those with mpg below
10th percentile: {len(filtered_mpg)}")

```

Project 4 Output

#4-1:

	mpg	cylinders	displacement	horsepower	...
387	27.0	4	140.0	86.0	
388	44.0	4	97.0	52.0	
389	32.0	4	135.0	84.0	
390	28.0	4	120.0	79.0	
391	31.0	4	119.0	82.0	

#4-2:

Number of missing values replaced in 'horsepower': 6

#4-3:

Number of cars from 'usa' with 4 cylinders: 72

#4-4:

Car with the highest mpg:

mazda glc - MPG: 46.6

#4-5:

Car with the highest 'displacement_to_power' ratio:

peugeot 504 - Ratio: 2.5396825396825395

#4-6:

Origin with the highest average mpg:

europe - Average MPG: 27.89

#4-7:

Number of cars remaining after removing those

with mpg below 10th percentile: 356

PROJECT 5 SOLUTIONS

```
import seaborn as sns
import pandas as pd

# Loading the datasets
tips = sns.load_dataset('tips')
flights = sns.load_dataset('flights')

#5-1
concat_vertically = pd.concat([tips, flights], axis=0)
print("#5-1:")
print(concat_vertically.head())

#5-2
df_discount = pd.DataFrame({
    'total_bill': [16.99, 10.34, 21.01],
    'discount': [10, 5, 7]
})
merged_df = pd.merge(tips, df_discount, on='total_bill',
    how='inner')
print("\n#5-2:")
print(merged_df.head())

#5-3
df_feedback = pd.DataFrame({
    'total_bill': [16.99, 10.34, 21.01],
    'feedback': ['Good', 'Average', 'Excellent']
})
left_joined = pd.merge(tips, df_feedback, on='total_bill', how='left')
print("\n#5-3:")
print(left_joined.head())

#5-4
```

```

df_review = pd.DataFrame({
    'total_bill': [16.99, 10.34, 15.00],
    'restaurant_review': ['4 stars', '3 stars', '5 stars']
})
outer_joined = pd.merge(tips, df_review, on='total_bill', how='outer')
print("\n#5-4:")
print(outer_joined.head())

#5-5
concat_horizontally = pd.concat([tips, flights], axis=1)
print("\n#5-5:")
print(concat_horizontally.head())

#5-6
df_experience = pd.DataFrame({
    'total_bill': [16.99, 10.34, 21.01],
    'tip': [1.01, 1.66, 3.50],
    'dining_experience': ['Casual', 'Formal', 'Casual']
})
merged_overlapping = pd.merge(tips, df_experience, on=['total_bill', 'tip'],
how='inner')
print("\n#5-6:")
print(merged_overlapping.head())

#5-7
multi_index_df = pd.DataFrame({
    'total_bill': [16.99, 10.34, 21.01],
    'tip': [1.01, 1.66, 3.50],
    'meal_preference': ['Veg', 'Non-Veg', 'Non-Veg']
}).set_index(['total_bill', 'tip'])
merged_multiindex = pd.merge(tips.set_index(['total_bill', 'tip']),
multi_index_df, left_index=True, right_index=True)
print("\n#5-7:")
print(merged_multiindex.head())

```

Project 5 Output

#5-1:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2.0	
1	10.34	1.66	Male	No	Sun	Dinner	3.0	
2	21.01	3.50	Male	No	Sun	Dinner	3.0	
3	23.68	3.31	Male	No	Sun	Dinner	2.0	
4	24.59	3.61	Female	No	Sun	Dinner	4.0	

#5-2:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2	
1	10.34	1.66	Male	No	Sun	Dinner	3	
2	21.01	3.50	Male	No	Sun	Dinner	3	

#5-3:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2	
1	10.34	1.66	Male	No	Sun	Dinner	3	

#5-4:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2	
1	10.34	1.66	Male	No	Sun	Dinner	3	

#5-5:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2.0	
1	10.34	1.66	Male	No	Sun	Dinner	3.0	
2	21.01	3.50	Male	No	Sun	Dinner	3.0	
3	23.68	3.31	Male	No	Sun	Dinner	2.0	
4	24.59	3.61	Female	No	Sun	Dinner	4.0	

#5-6:

	total_bill	tip	sex	smoker	day	time	size	...
0	16.99	1.01	Female	No	Sun	Dinner	2	
1	10.34	1.66	Male	No	Sun	Dinner	3	

#5-7:

sex	smoker	day	time	size	...
-----	--------	-----	------	------	-----

total_bill tip

16.99 1.01 Female No Sun Dinner 2

10.34 1.66 Male No Sun Dinner 3

PROJECT 6 SOLUTIONS

```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import scipy.stats as stats

# Load the dataset
tips = sns.load_dataset('tips')

#6-1:
tips['total_bill'].hist(edgecolor='black')
plt.xlabel('Total Bill')
plt.ylabel('Frequency')
plt.title('Histogram of Total Bill')
plt.show()

#6-2:
correlation_matrix = tips.corr()
print(correlation_matrix)

#6-3:
tips.plot.scatter(x='total_bill', y='tip')
plt.xlabel('Total Bill')
plt.ylabel('Tip')
plt.title('Scatter plot of Total Bill vs Tip')
plt.show()

#6-4:
skewness = tips['total_bill'].skew()
kurtosis = tips['total_bill'].kurt()
print(f"Skewness: {skewness}")
print(f"Kurtosis: {kurtosis}")

#6-5:
```

```
value_counts = tips['day'].value_counts()
print(value_counts)
value_counts.plot(kind='bar')
plt.xlabel('Day')
plt.ylabel('Frequency')
plt.title('Frequency of unique values in Day column')
plt.show()
```

#6-6:


```
tips.boxplot(column='total_bill')
plt.xlabel('Total Bill')
plt.ylabel('Value')
plt.title('Box plot of Total Bill')
plt.show()
```

#6-7:

```
Q1 = tips['total_bill'].quantile(0.25)
Q3 = tips['total_bill'].quantile(0.75)
IQR = Q3 - Q1
print(f"IQR for total_bill: {IQR}")
```

Project 6 Output


#6-1:



#6-2:

```
total_bill      tip      size
total_bill  1.000000  0.675734  0.598315
tip        0.675734  1.000000  0.489299
size       0.598315  0.489299  1.000000
```

#6-3:




#6-4:

Skewness: 1.1332130376158205
Kurtosis: 1.2184840156638854


#6-5:

```
Sat    87  
Sun    76  
Thur   62  
Fri    19  
Name: day, dtype: int64  
# Displays bar chart with days
```

#6-5:



#6-6:



#6-7:

IQR for total_bill: 10.77999999999998

PROJECT 7 SOLUTIONS

```
import seaborn as sns
import numpy as np
import pandas as pd

# Load the dataset
diamonds = sns.load_dataset('diamonds')

#7-1:
filtered_diamonds = diamonds.query('cut == "Ideal" and price > 500')
print(filtered_diamonds)

#7-2:
pattern_matched_rows =
    diamonds[diamonds['color'].str.match('^[EF]$')]
print(pattern_matched_rows)

#7-3:
diamonds_with_color_index = diamonds.set_index('color')
reset_diamonds = diamonds_with_color_index.reset_index()
print(reset_diamonds)

#7-4:
diamonds['price_per_carat'] = diamonds.eval('price / carat')
print(diamonds[['price', 'carat', 'price_per_carat']])

#7-5:
diamonds['high_price'] = np.where(diamonds['price'] > 1000, True, False)
print(diamonds[['price', 'high_price']])

#7-6:
low_variance_cols =
    diamonds.var()[diamonds.var() < 1].index.tolist()
diamonds_dropped = diamonds.drop(columns=low_variance_cols)
```

```
print(diamonds_dropped)
```

#7-7:

```
melted_diamonds = diamonds.melt(id_vars=['color', 'cut'])  
print(melted_diamonds)
```

Project 7 Output

#7-1:

	carat	cut	color	clarity	depth	table	price	...
60	0.35	Ideal	I	VS1	60.9	57.0	552	
62	0.30	Ideal	D	SI1	62.5	57.0	552	
63	0.30	Ideal	D	SI1	62.1	56.0	552	
65	0.28	Ideal	G	VVS2	61.4	56.0	553	
66	0.32	Ideal	I	VVS1	62.0	55.3	553	
...	
53925	0.79	Ideal	I	SI1	61.6	56.0	2756	
53926	0.71	Ideal	E	SI1	61.9	56.0	2756	
53929	0.71	Ideal	G	VS1	61.4	56.0	2756	
53935	0.72	Ideal	D	SI1	60.8	57.0	2757	
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	

[20919 rows x 10 columns]

#7-2

	carat	cut	color	clarity	depth	table	price	...
0	0.23	Ideal	E	SI2	61.5	55.0	326	
1	0.21	Premium	E	SI1	59.8	61.0	326	
2	0.23	Good	E	VS1	56.9	65.0	327	
8	0.22	Fair	E	VVS2	65.1	61.0	337	
12	0.22	Premium	F	SI1	60.4	61.0	342	
...	
53928	0.79	Premium	E	SI2	61.4	58.0	2756	
53930	0.71	Premium	E	SI1	60.5	55.0	2756	
53931	0.71	Premium	F	SI1	59.8	62.0	2756	
53932	0.70	Very Good	E	VVS2	60.5	59.0	2757	

53933 0.70 Very Good E VS2 61.2 59.0 2757

[19339 rows x 10 columns]

#7-3

	color	carat	cut	clarity	depth	table	price	...
0	E	0.23	Ideal	SI2	61.5	55.0	326	
1	E	0.21	Premium	SI1	59.8	61.0	326	
2	E	0.23	Good	VS1	56.9	65.0	327	
3	I	0.29	Premium	VS2	62.4	58.0	334	
4	J	0.31	Good	SI2	63.3	58.0	335	
...
53935	D	0.72	Ideal	SI1	60.8	57.0	2757	
53936	D	0.72	Good	SI1	63.1	55.0	2757	
53937	D	0.70	Very Good	SI1	62.8	60.0	2757	
53938	H	0.86	Premium	SI2	61.0	58.0	2757	
53939	D	0.75	Ideal	SI2	62.2	55.0	2757	

[53940 rows x 10 columns]

#7-4

	price	carat	price_per_carat
0	326	0.23	1417.391304
1	326	0.21	1552.380952
2	327	0.23	1421.739130
3	334	0.29	1151.724138
4	335	0.31	1080.645161
...
53935	2757	0.72	3829.166667
53936	2757	0.72	3829.166667
53937	2757	0.70	3938.571429
53938	2757	0.86	3205.813953
53939	2757	0.75	3676.000000

[53940 rows x 3 columns]

#7-5

price	high_price
-------	------------

```

0      326    False
1      326    False
2      327    False
3      334    False
4      335    False
...
...   ...
53935 2757    True
53936 2757    True
53937 2757    True
53938 2757    True
53939 2757    True

```

[53940 rows x 2 columns]

#7-6

		cut	color	clarity	depth	table	price	...
0		Ideal	E	SI2	61.5	55.0	326	
1		Premium	E	SI1	59.8	61.0	326	
2		Good	E	VS1	56.9	65.0	327	
3		Premium	I	VS2	62.4	58.0	334	
4		Good	J	SI2	63.3	58.0	335	
...
53935		Ideal	D	SI1	60.8	57.0	2757	
53936		Good	D	SI1	63.1	55.0	2757	
53937		Very Good	D	SI1	62.8	60.0	2757	
53938		Premium	H	SI2	61.0	58.0	2757	
53939		Ideal	D	SI2	62.2	55.0	2757	

#7-7

	color	cut	variable	value
0	E	Ideal	carat	0.23
1	E	Premium	carat	0.21
2	E	Good	carat	0.23
3	I	Premium	carat	0.29
4	J	Good	carat	0.31
...
539395	D	Ideal	high_price	True

539396	D	Good	high_price	True
539397	D	Very Good	high_price	True
539398	H	Premium	high_price	True
539399	D	Ideal	high_price	True

[539400 rows x 4 columns]

PROJECT 8 SOLUTIONS

```
import seaborn as sns
import pandas as pd
from sklearn.preprocessing import StandardScaler,
OneHotEncoder
from sklearn.decomposition import PCA
from statsmodels.stats.outliers_influence
import variance_inflation_factor

# Load dataset
titanic = sns.load_dataset('titanic')

# 8-1
titanic['family_size'] = titanic['sibsp'] + titanic['parch']
print("#8-1:\n", titanic['family_size'].head())

# 8-2
encoder = OneHotEncoder(drop='first', sparse=False)
encoded_class = encoder.fit_transform(titanic[['class']])
encoded_df = pd.DataFrame(encoded_class,
columns=encoder.get_feature_names(['class']))
titanic = pd.concat([titanic, encoded_df], axis=1)
print("#8-2:\n", titanic[encoded_df.columns].head())

# 8-3
scaler = StandardScaler()
titanic['age_z'] = scaler.fit_transform(titanic[['age']])
print("#8-3:\n", titanic['age_z'].head())

# 8-4
# Note: Normally, you'd remove categorical variables or one-hot encode them, but for simplicity, we'll use only 'age' and 'fare' here.
pca = PCA(n_components=1)
```

```
titanic['pca_component'] = pca.fit_transform(titanic[['age', 'fare']])
print("#8-4:\n", titanic['pca_component'].head())

# 8-5
# Computing VIF for age, fare, and the one-hot encoded class columns
features = ['age', 'fare', 'class_Second', 'class_Third']
vif_data = pd.DataFrame()
vif_data['feature'] = features
vif_data['VIF'] = [variance_inflation_factor(titanic[features].values, i) for i
in range(titanic[features].shape[1])]
print("#8-5:\n", vif_data)
```

```
# 8-6
bins = [0, 18, 35, 60, 100]
labels = ['child', 'young_adult', 'adult', 'senior']
titanic['age_group'] = pd.cut(titanic['age'], bins=bins, labels=labels,
right=False)
print("#8-6:\n", titanic['age_group'].head())
```

```
# 8-7
mean_encode = titanic.groupby('embark_town')['survived'].mean()
titanic['embark_encoded'] =
    titanic['embark_town'].map(mean_encode)
print("#8-7:\n", titanic['embark_encoded'].head())
```

Project 8 Output

```
#8-1:
0 1
1 1
2 0
3 1
4 0
Name: family_size, dtype: int64
```

```
#8-2:
class_Second  class_Third
```

```
0      0.0      1.0
1      0.0      0.0
2      0.0      1.0
3      0.0      0.0
4      0.0      1.0
```

#8-3:

```
0 -0.592481
1 0.638789
2 -0.284663
3 0.407926
4 0.407926
```

Name: age_z, dtype: float64

#8-4:

```
0 -20.493525
1 40.918998
2 -18.387374
3 23.506727
4 -17.887235
```

Name: pca_component, dtype: float64

#8-5:

	feature	VIF
0	age	4.395834
1	fare	1.691754
2	class_Second	1.859111
3	class_Third	2.931346

#8-6:

```
0 young_adult
1      adult
2 young_adult
3      adult
4      adult
```

Name: age_group, dtype: category

Categories (4, object): ['child' < 'young_adult' < 'adult' < 'senior']

#8-7:

0